

A close-up photograph of rowers in a boat, showing their hands on yellow handles and black oars. The oars are dipping into the water, creating ripples. The rowers are wearing blue long-sleeved shirts.

CSC353

# More about Synchronization

Download class materials from  
[university.xamarin.com](https://university.xamarin.com)



**Xamarin** University

Information in this document is subject to change without notice. The example companies, organizations, products, people, and events depicted herein are fictitious. No association with any real company, organization, product, person or event is intended or should be inferred. Complying with all applicable copyright laws is the responsibility of the user.

Microsoft or Xamarin may have patents, patent applications, trademarks, copyrights, or other intellectual property rights covering subject matter in this document. Except as expressly provided in any license agreement from Microsoft or Xamarin, the furnishing of this document does not give you any license to these patents, trademarks, or other intellectual property.

© 2014-2017 Xamarin Inc., Microsoft. All rights reserved.


Xamarin, MonoTouch, MonoDroid, Xamarin.iOS, Xamarin.Android, Xamarin Studio, and Visual Studio are either registered trademarks or trademarks of Microsoft in the U.S.A. and/or other countries.

Other product and company names herein may be the trademarks of their respective owners.

# Objectives

1. Beyond Monitor – implement other synchronization strategies
2. Work with Synchronized Collections





# Implement other synchronization strategies


# Tasks

1. Examine Synchronization strategies
2. Use Wait Handles to control execution of multiple threads
3. Control multiple thread entry to a resource
4. Coordinate thread activity




# Synchronization strategies

- ❖ Monitors are the primary strategy to synchronize data access, but .NET provides other synchronization mechanisms for specific cases



Efficient waiting



High reader vs.  
writer ratio



Allow  $n$  accesses  
simultaneously



Coordinate activity

# Reminder: multiple lock coordination

- ❖ Be careful about waiting on multiple monitors – this will almost always produce race conditions and deadlock scenarios in your code

```
public class Account
{
    object accountLock = new object();

    public static void Transfer(Account accountFrom, Account accountTo, decimal amount)
    {
        lock (accountFrom.accountLock)
        {
            lock (accountTo.accountLock)
            {
                ... // Transfer money
            }
        }
    }
}
```

What if some other thread owns this lock and is waiting on the first one?



# Locking with Mutex

- ❖ A **Mutex** is a locking object that can be used across processes – although this extra power is not typically useful in mobile applications

```
public class Account
{
    Mutex accountLock = new Mutex();

    public static void Transfer(Account accountFrom, Account accountTo, decimal amount)
    {
        accountFrom.accountLock.WaitOne();
        accountTo.accountLock.WaitOne();

        ... // Transfer money

        accountFrom.accountLock.ReleaseMutex();
        accountTo.accountLock.ReleaseMutex();
    }
}
```

WaitOne methods  
acquire the lock

The ReleaseMutex  
method releases the lock



# WaitHandle coordination

- ❖ **Mutex** derives from **WaitHandle** so it can be used with the **WaitHandle** static synchronization methods such as **WaitAll** and **WaitAny**

```
public class Account
{
    Mutex accountLock = new Mutex();

    public static void Transfer(Account accountFrom, Account accountTo, decimal amount)
    {
        WaitHandle.WaitAll(new[] { accountFrom.accountLock, accountTo.accountLock });
        try {
            ... // Transfer money
        }
        finally {
            accountFrom.accountLock.ReleaseMutex();
            accountTo.accountLock.ReleaseMutex();
        }
    }
}
```

The operating system ensures that no deadlock will occur here and that this thread will own both mutexes before returning

# Release Mutex

- ❖ Always call **ReleaseMutex** when you are finished – this releases the lock and allows other clients to acquire the lock and access the resource

```
public class Account
{
    Mutex accountLock = new Mutex();

    public static void Transfer(Account accountFrom, Account accountTo, decimal amount)
    {
        WaitHandle.WaitAll(new[] { accountFrom.accountLock, accountTo.accountLock });
        try {
            ... // Transfer money
        }
        finally {
            accountFrom.accountLock.ReleaseMutex();
            accountTo.accountLock.ReleaseMutex();
        }
    }
}
```

Must always release the mutexes when you are finished accessing the data


# Reader vs. writer ratio

- ❖ Monitors block all thread access to a block of code while the monitor is owned; even if all the thread wants to do is *read* the data
- ❖ In these cases, it would be safe to allow multiple threads to access the resource as long as they do not modify any data



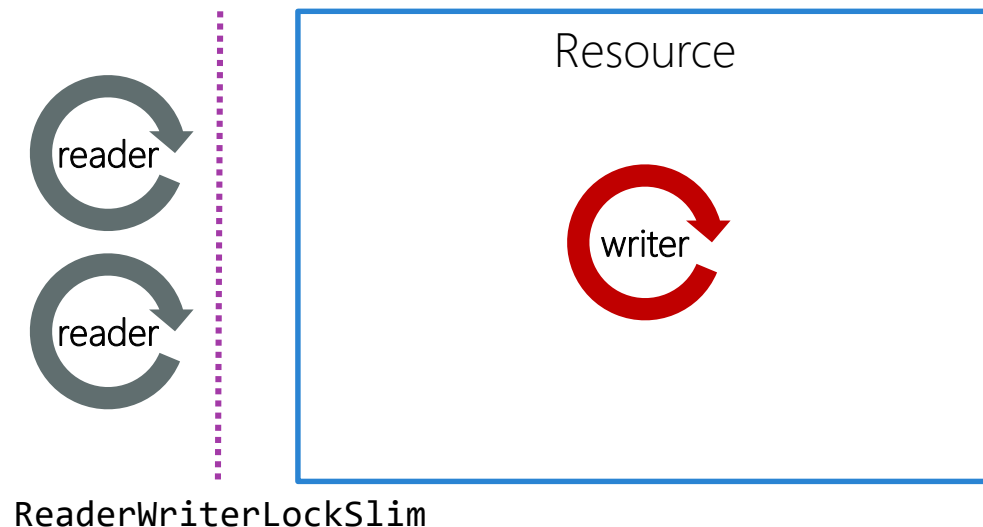
# Introducing ReaderWriterLockSlim

- ❖ **ReaderWriterLockSlim** is a locking mechanism that allows multiple readers simultaneous access to a resource, but only a single writer at a time



# Introducing ReaderWriterLockSlim

- ❖ **ReaderWriterLockSlim** is a locking mechanism that allows multiple readers simultaneous access to a resource, but only a single writer at a time



# Using ReaderWriterLockSlim


```
Dictionary<string, string> cache = new Dictionary<string, string>();  
ReaderWriterLockSlim rwls = new ReaderWriterLockSlim();
```

```
string GetCachedValue(string key) {  
    rwls.EnterReadLock();  
    try { return cache[key]; }  
    finally { rwls.ExitReadLock(); }  
}
```

Multiple threads are allowed to enter and exit the read lock simultaneously

```
void SetCacheValue(string key, string value) {  
    rwls.EnterWriteLock();  
    try { cache.Add(key, value); }  
    finally { rwls.ExitWriteLock(); }  
}
```

.. but only one thread (at a time) can enter the write lock, all read lock requests are blocked while the write lock is owned

 Be careful: There are several older classes with similar names that are either deprecated or more expensive to use – should generally prefer types that end with **Slim**

# Should I prefer RWLS?

- ❖ Always start with Monitor; it is, by far, the most efficient lock in .NET
- ❖ For infrequent update cases, try replacing with RWLS *and then profile the application*
- ❖ For certain scenarios it may be more efficient, for others it might not be; must test it to find out





# Demonstration

Using ReaderWriterLockSlim



**Xamarin**  
University


# Allowing multiple threads

- ❖ Monitors allow one thread at a time; but sometimes we'd like to allow for more than one thread, up to a known limit
- ❖ This is done with a *semaphore* which allows **N** number of simultaneous accesses
- ❖ New threads cannot acquire semaphore until an existing owner releases it



# Introducing SemaphoreSlim

- ❖ **SemaphoreSlim** allows for a known number of simultaneous acquisitions



# Using SemaphoreSlim

```
public class LicenseServer
{
    SemaphoreSlim guard = new SemaphoreSlim(3); // 3 simultaneous requests allowed

    public async Task<string> RequestLicense(string key) {
        await guard.WaitAsync();

        try {
            return await InternalRequest(key);
        }
        finally {
            guard.Release();
        }
    }
    ...
}
```

Initialize with the allowed count

# Using SemaphoreSlim

```
public class LicenseServer
{
    SemaphoreSlim guard = new SemaphoreSlim(3); // 3 simultaneous requests allowed

    public async Task<string> RequestLicense(string licenseKey)
    {
        await guard.WaitAsync();

        try {
            return await InternalRequestLicense(licenseKey);
        }
        finally {
            guard.Release();
        }
    }
    ...
}
```

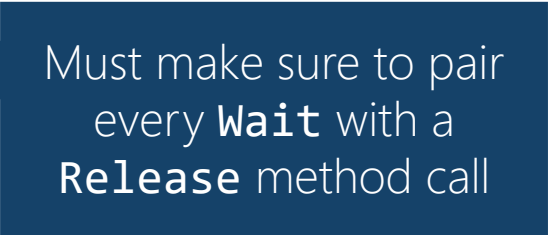
Attempt to acquire the semaphore – class includes both **async** and blocking versions of **Wait** method; be aware that the lock is *not* reentrant, even for the same thread

# Using SemaphoreSlim


```
public class LicenseServer
{
    SemaphoreSlim guard = new SemaphoreSlim(3); // 3 simultaneous requests allowed

    public async Task<string> RequestLicenseCodeAsync(string key) {
        await guard.WaitAsync();

        try {
            return await InternalRequest(key);
        }
        finally {
            guard.Release();
        }
    }
    ...
}
```



Must make sure to pair every **Wait** with a **Release** method call



There are several variations of **Wait** available, including ones that take **CancellationToken**s to be able to cancel the wait from an external request

# Demonstration

Using Semaphore



**Xamarin**  
University



# Coordinating activity

- ❖ Sometimes the synchronization required is to *wait* for an event to occur to coordinate two or more threads



Something  
happened




Countdown  
coordination



Rendezvous  
coordination

# Coordination terminology

- ❖ Several common terms / methods are used when using coordination objects



# Coordination terminology

- ❖ Several common terms / methods are used when using coordination objects



Wait

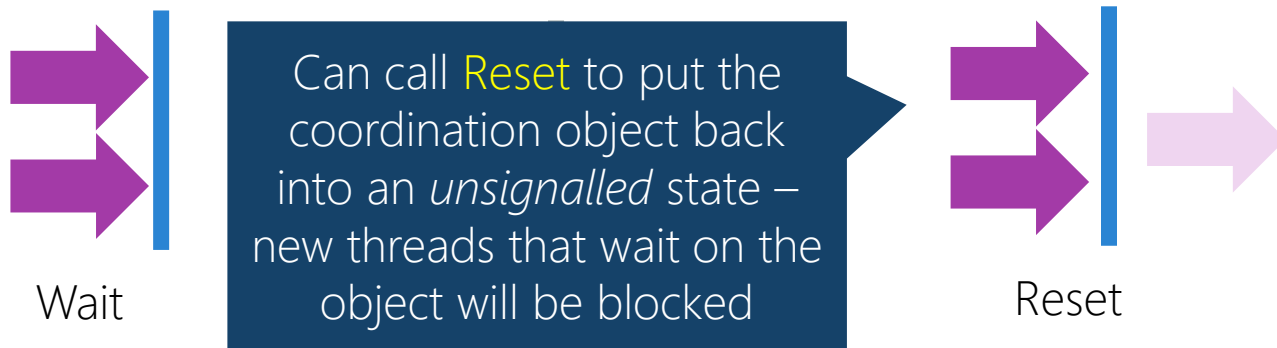


Signal or Set

When the event *does* happen, a thread calls **Signal** (or sometimes **Set**) to tell any waiters


# Coordination terminology

- ❖ Several common terms / methods are used when using coordination objects



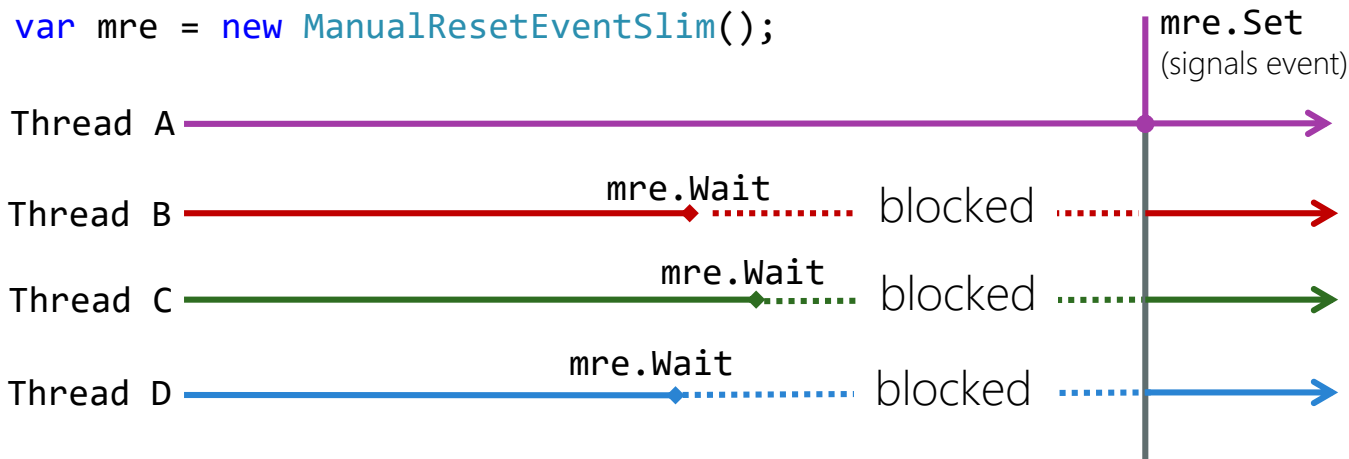
# Event-based coordination

- ❖ Sometimes, two or more background threads need to coordinate on a specific action occurring
- ❖ This can be achieved using a **ManualResetEventSlim** object
- ❖ Particularly useful when *multiple* threads must be coordinated to a specific action



# Event-based coordination

- ❖ **ManualResetEventSlim** releases all waiting threads when it is signaled; it is used to indicate a specific action has occurred in a thread



# Using ManualResetEventSlim

- ❖ We want to build a thread-safe, blocking queue

```
BlockingQueue<string> queue = new BlockingQueue<string>();
```

```
string newValue = queue.Dequeue();
```

```
queue.Enqueue
```

Multiple threads can call **Dequeue**, each will block until data is available, and only one thread will see the data



# Using ManualResetEventSlim

- ❖ We want to build a thread-safe, blocking queue

```
BlockingQueue<string> queue = new BlockingQueue<string>();
```

```
string newValue = queue.Dequeue
```

Multiple threads can call **Enqueue**, and if threads are waiting on data, one will process it

```
queue.Enqueue("Pass to a thread!");
```

# Event-based coordination

- ❖ Blocking queue will utilize a `Queue<T>` and a `ManualResetEventSlim` to provide notification about data availability

```
public class BlockingQueue<T>
{
    Queue<T> queue = new Queue<T>();
    ManualResetEventSlim hasData = new ManualResetEventSlim(false);
    ...
}
```

Event starts in *un-signaled* state by default, can specify initial state in constructor or call **Reset** to change it back to un-signaled later

# Event-based coordination

- ❖ **Enqueue** method will lock queue, add our item and then signal the event if this is the only data element available to wake up any waiters

```
public class BlockingQueue<T>
{
    public void Enqueue(T obj) {
        lock (queue) {
            queue.Enqueue(obj);
            if (queue.Count == 1)
                hasData.Set();
        }
    }
    ...
}
```

To unblock waiting threads, must signal the event by calling **Set**; any waiters are then made ready for scheduling and have the opportunity to run

# Event-based coordination

- ❖ **Dequeue** will lock queue and then release the lock and wait if no data is available; once data arrives, we lock queue and try to get the data

```
public T Dequeue() {  
    T value;  
    bool taken = true; Monitor.Enter(queue);  
    try {  
        while (queue.Count == 0) {  
            taken = false; Monitor.Exit(queue);  
            hasData.Wait();  
            Monitor.Enter(queue); taken = true;  
        }  
        value = queue.Dequeue();  
        if (queue.Count == 0)  
            hasData.Reset();  
    }  
    finally { if (taken) Monitor.Exit(queue); }  
    return value;  
}
```

Wait causes thread to *block* until event is signaled; can pass optional timeout, or **CancellationToken** to provide for cancellation

# Event-based coordination

- ❖ **Dequeue** must also reset the event once we have no more data available

```
public T Dequeue() {  
    T value;  
    bool taken = true; Monitor.Enter(queue);  
    try {  
        while (queue.Count == 0) {  
            taken = false; Monitor.Exit(queue);  
            hasData.Wait();  
            Monitor.Enter(queue); taken = true;  
        }  
        value = queue.Dequeue();  
        if (queue.Count == 0)  
            hasData.Reset();  
    }  
    finally { if (taken) Monitor.Exit(queue); }  
    return value;  
}
```

**Reset** causes the event to move back to the *unsigaled* state so that new threads that call **Wait** will be blocked on the event

# Countdown coordination

- ❖ A common coordination style is to wait until a set number of activities have occurred or completed using a **CountdownEvent**

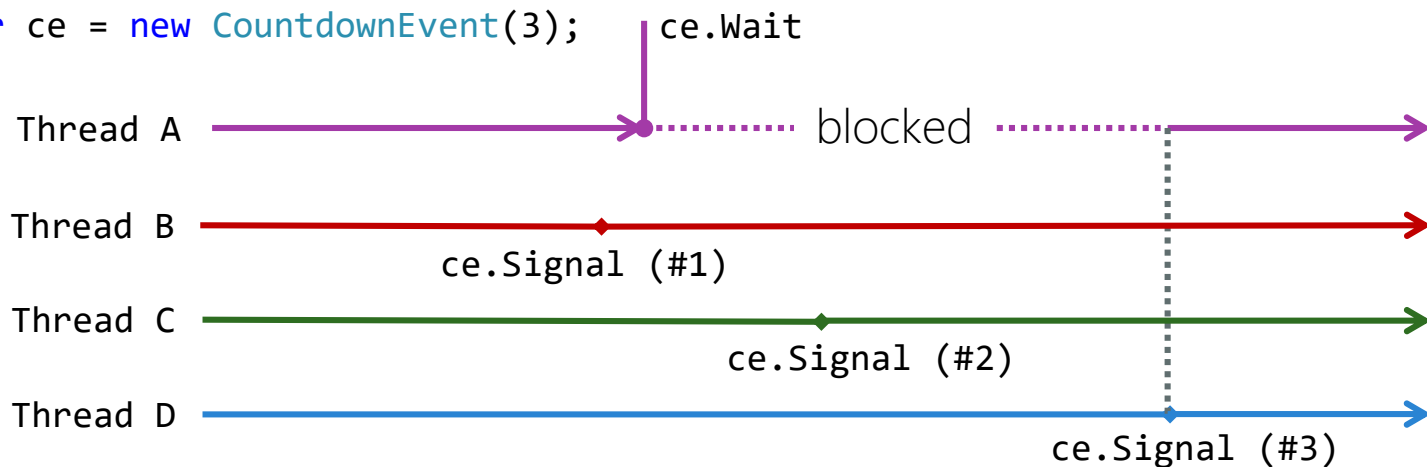


5

# Countdown coordination

- ❖ **CountdownEvent** unblocks waiters when it has been signaled a specific number of times

```
var ce = new CountdownEvent(3);
```





# Countdown coordination

- ❖ **CountdownEvent** allows one thread to wait for **N** operations to occur

```
CountdownEvent cde = new CountdownEvent(5);  
for (int i = 0; i < cde.InitialCount; i++)  
{  
    Task.Run(() => {  
        // Do work  
        cde.Signal();  
    });  
}  
  
cde.Wait(1000);  
int doneAfter1Second = cde.InitialCount - cde.CurrentCount;  
cde.Wait();
```

Must indicate how many times we will be signaled

# Countdown coordination

- ❖ **CountdownEvent** allows one thread to wait for **N** operations to occur

```
CountdownEvent cde = new CountdownEvent(5);
for (int i = 0; i < cde.InitialCount; i++)
{
    Task.Run(() => {
        // Do work
        cde.Signal();
    });
}

cde.Wait(1000);
int doneAfter1Second = cde.InitialCount - cde.CurrentCount;
cde.Wait();
```

Each subordinate task/thread then signals the **CountdownEvent**

# Countdown coordination

- ❖ **CountdownEvent** allows one thread to wait for **N** operations to occur

```
CountdownEvent cde = new CountdownEvent(5);
for (int i = 0; i < cde.InitialCount; i++)
{
    Task.Run(() => {
        // Do work
        cde.Signal();
    });
}

cde.Wait(1000);
int doneAfter1Second = cde.InitialCount - cde.CurrentCount;
cde.Wait();
```

Can block waiting for all the operations to occur; supplying an optional timeout or **CancellationToken**

# Countdown coordination

❖ **CountdownEvent** allows one thread to wait for **N** operations to occur

```
CountdownEvent cde = new CountdownEvent(5);  
for (int i = 0; i < cde.InitialCount; i++)  
{  
    Task.Run(() => {  
        // Do work  
        cde.Signal();  
    });  
}
```

Can look at properties to  
determine how many  
operations have been  
counted so far

```
cde.Wait(1000);  
int doneAfter1Second = cde.InitialCount - cde.CurrentCount;  
cde.Wait();
```

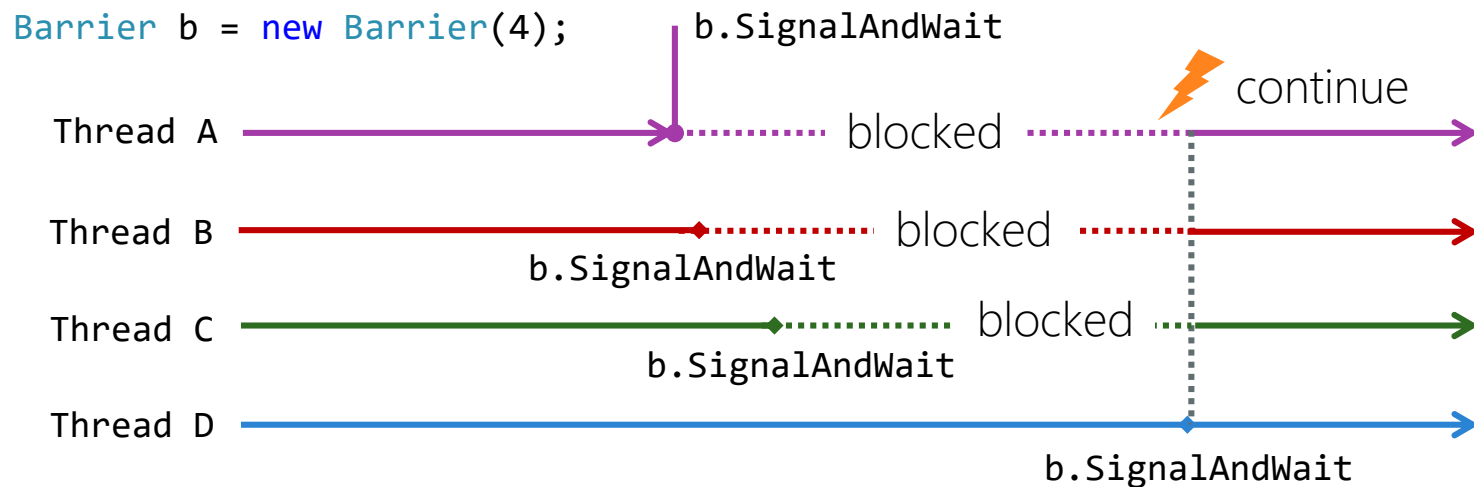
# Rendezvous coordination

- ❖ Sometimes useful to know when a set of threads have hit a certain point in their processing
- ❖ Then have each thread *wait* until all other threads have also hit that point; once all are in phase, we let the threads continue
- ❖ The **Barrier** class allows for this sort of "rendezvous" coordination



# Rendezvous coordination

- ❖ **Barrier** class provides for "rendezvous" style coordination where our work is done in *phases* and must be synchronized between each phase



# Rendezvous coordination

- ❖ **Barrier** class provides for "rendezvous" style coordination where our work is done in *phases* and must be synchronized between each phase

```
Barrier b = new Barrier(3);  
void RunAHorse() {  
    ... // Initialize the horse  
    // Wait on the gate (barrier)  
    b.SignalAndWait();  
    // All horses ready, gate is open - GO!  
}  
...  
for (int i = 0; i < 3; i++) {  
    Task.Run(RunAHorse);  
}
```

Must specify how many signals to require before unblocking callers

# Rendezvous coordination

- ❖ **Barrier** class provides for "rendezvous" style coordination where our work is done in *phases* and must be synchronized between each phase

```
Barrier b = new Barrier(3);  
void RunAHorse() {  
    ... // Initialize the horse  
    // Wait on the gate (barrier)  
    b.SignalAndWait();  
    // All horses ready, gate opens  
}  
...  
for (int i = 0; i < 3; i++) {  
    Task.Run(RunAHorse);  
}
```

Threads will block until  
**SignalAndWait** is  
called 3 times



# Rendezvous coordination

- ❖ **Barrier** class provides for "rendezvous" style coordination where our work is done in *phases* and must be synchronized between each phase

```
Barrier b = new Barrier(3);  
void RunAHorse() {  
    ... // Initialize the horse  
    // Wait on the gate (barrier)  
    b.SignalAndWait();  
    // All horses ready, gate is open - GO!  
}  
...  
for (int i = 0; i < 3; i++) {  
    Task.Run(RunAHorse);  
}
```

Once we hit 3 calls to **SignalAndWait**, all threads are released and the barrier should be disposed

# Dealing with dynamic participants

- ❖ Can add and remove additional counts from the **Barrier** at runtime

```
Barrier b = new Barrier(1);  
...  
b.AddParticipant(); // Now 2  
...  
b.RemoveParticipant(); // Back to one
```

# Barrier Phasing

- ❖ Can supply a *phase action* to be executed after all threads have reached the barrier, but *before* the threads are released to run

```
Barrier barrier = new Barrier(NumThreads, new Action(VerifyBalances));

void DoTransfers(List<TransferRequest> requests) {
    ... // Transfer money
    barrier.SignalAndWait();
    ...
}

void VerifyBalances() { ... }
```

Method is called once all threads are blocked on the barrier

# Summary

1. Examine Synchronization strategies
2. Use Wait Handles to control execution of multiple threads
3. Control multiple thread entry to a resource
4. Coordinate thread activity



# Working with Synchronized Collections

# Tasks

1. Explore thread-safe collections
2. Use concurrent collections
3. Understand the  
Producer/Consumer pattern



# Collections in .NET

- ❖ Standard collections in .NET are not thread safe and must always be synchronized in your code if they are accessed in a read/write fashion by multiple threads

```
object collectionGuard = new object();
List<string> data = new List<string>();

public void Add(string text)
{
    lock (collectionGuard) {
        data.Add(text);
    }
}
```

# Thread-safe collections

- ❖ .NET 4.5 introduced a set of collections in the namespace **System.Collections.Concurrent** which ensure thread-safety; either by using locks internally, or through the use of lock-free, thread-safe algorithms

`ConcurrentQueue<T>`

Provides a first-in-first-out (FIFO) collection of data



# Thread-safe collections

- ❖ .NET 4.5 introduced a set of collections in the namespace **System.Collections.Concurrent** which ensure thread-safety; either by using locks internally, or through the use of lock-free, thread-safe algorithms

`ConcurrentQueue<T>`

`ConcurrentStack<T>`

Provides a last-in-first-out (LIFO) collection of data

# Thread-safe collections

- ❖ .NET 4.5 introduced a set of collections in the namespace **System.Collections.Concurrent** which ensure thread-safety; either by using locks internally, or through the use of lock-free, thread-safe algorithms

`ConcurrentQueue<T>`

`ConcurrentStack<T>`

`ConcurrentDictionary  
<K,V>`

Provides a collection  
of random-access  
key-value pairs

# Thread-safe collections

- ❖ .NET 4.5 introduced a set of collections in the namespace **System.Collections.Concurrent** which ensure thread-safety; either by using locks internally, or through the use of lock-free, thread-safe algorithms

**ConcurrentQueue<T>**

**ConcurrentStack<T>**

**Conc**

Provides an unordered  
collection of objects

**ConcurrentBag<T>**

# Using the concurrent collections

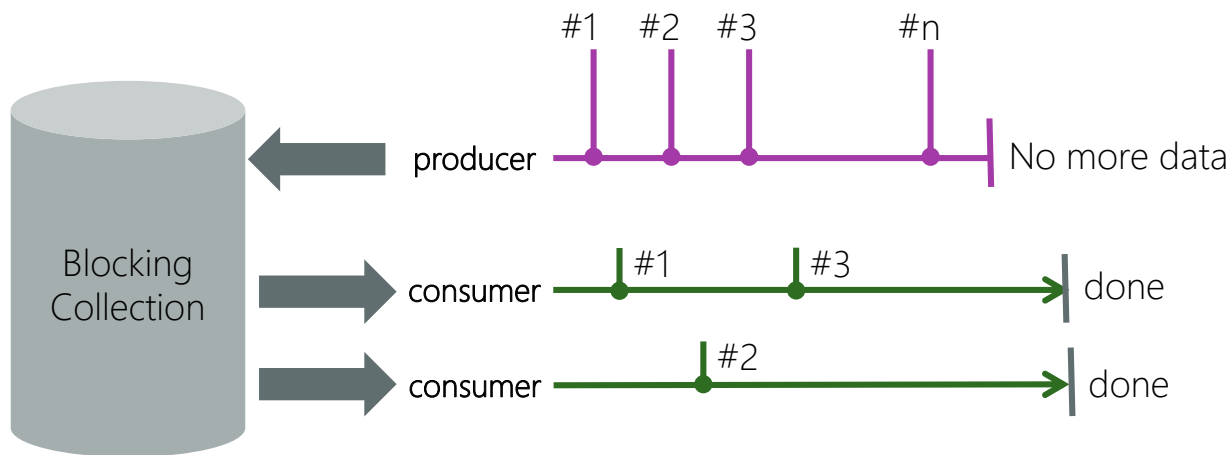
- ❖ Because multiple threads might be changing the collection, all retrieval operations are "non-deterministic"

```
ConcurrentQueue<string> data = new ConcurrentQueue<string>();  
  
public string GetNextItem()  
{  
    string rc = null;  
    return (data.TryDequeue(out rc)) ? rc : null;  
}
```

No need to take a lock,  
just try to dequeue value

# Producer / Consumer pattern

- ❖ New collections also provide support for **producer / consumer** pattern through a wrapper **BlockingCollection<T>**



# Producer / Consumer pattern

- ❖ **BlockingCollection** provides **IEnumerable** which automatically terminates when the producer indicates no more values are forthcoming

```
ConcurrentStack<int> stack = new ConcurrentStack<int>();  
BlockingCollection<int> work = new BlockingCollection<int>(stack);
```

```
void Producer() {  
    for (int i = 0; i < 100; i++)  
        work.Add(i);  
    work.CompleteAdding();  
}
```

```
void Consumer() {  
    foreach (int value in work.GetConsumingEnumerable()) {  
        // Process value  
    }  
}
```

Constructor takes the type of collection to use as the backing storage, defaults to queue

# Producer / Consumer pattern

- ❖ **BlockingCollection** provides **IEnumerable** which automatically terminates when the producer indicates no more values are forthcoming

```
ConcurrentStack<int> stack = new ConcurrentStack<int>();  
BlockingCollection<int> work = new BlockingCollection<int>(stack);  
  
void Producer() {  
    for (int i = 0; i < 100; i++)  
        work.Add(i);  
    work.CompleteAdding();  
}  
  
void Consumer() {  
    foreach (int value in work.GetConsumingEnumerable()) {  
        // Process value  
    }  
}
```

Producer adds values into  
the blocking collection

# Producer / Consumer pattern

- ❖ **BlockingCollection** provides **IEnumerable** which automatically terminates when the producer indicates no more values are forthcoming

```
ConcurrentStack<int> stack = new ConcurrentStack<int>();  
BlockingCollection<int> work = new BlockingCollection<int>(stack);  
  
void Producer() {  
    for (int i = 0; i < 100; i++)  
        work.Add(i);  
    work.CompleteAdding();  
}  
  
void Consumer() {  
    foreach (int value in work.GetConsumingEnumerable()) {  
        // Process value  
    }  
}
```

Consumer(s) read from  
blocking **IEnumerable**



# Producer / Consumer pattern

- ❖ **BlockingCollection** provides **IEnumerable** which automatically terminates when the producer indicates no more values are forthcoming

```
ConcurrentStack<int> stack = new ConcurrentStack<int>();  
BlockingCollection<int> work = new BlockingCollection<int>(stack);  
  
void Producer() {  
    for (int i = 0; i < 100;  
        work.Add(i);  
        work.CompleteAdding();  
}  
  
void Consumer() {  
    foreach (int value in work.GetConsumingEnumerable()) {  
        // Process value  
    }  
}
```

Producer then indicates when no more data is available, this ends the **IEnumerable** loop

# Demonstration

Blocking Collection Demo



**Xamarin**  
University

# Summary

1. Explore thread-safe collections
2. Use concurrent collections
3. Understand the  
Producer/Consumer pattern



# Thank You!

Please complete the class survey in your profile:  
[university.xamarin.com/profile](https://university.xamarin.com/profile)

